

# Performance Evaluation and Applications

 POLITECNICO DI MILANO

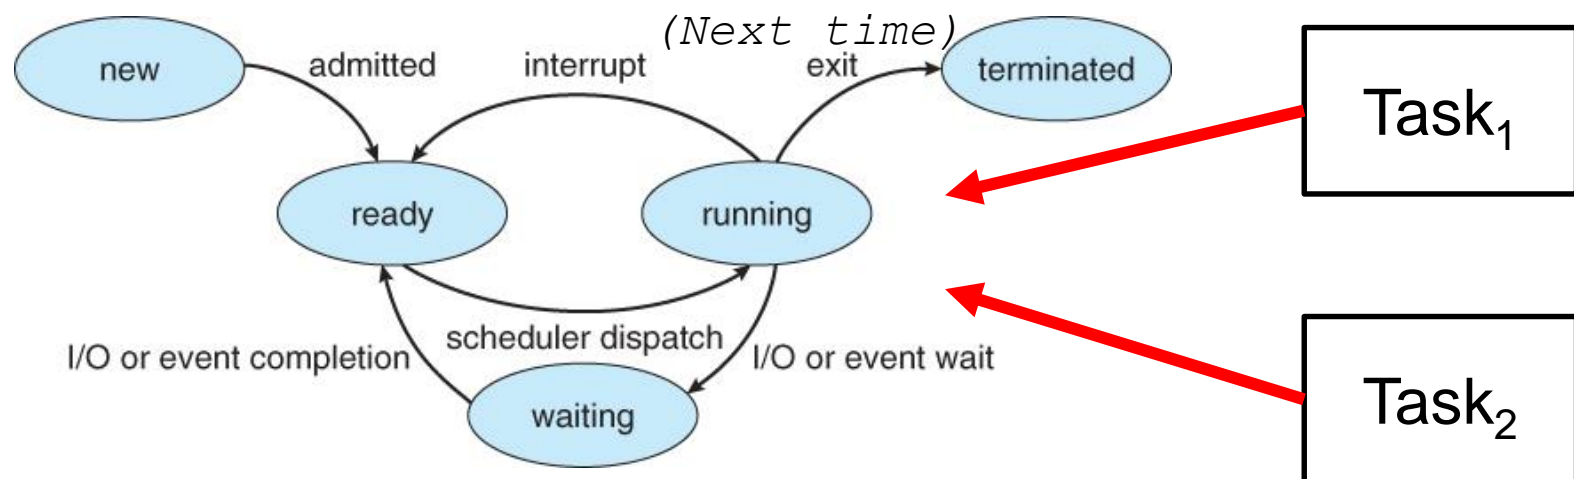
## Race, Concurrency, and the Exponential Distribution

POLITECNICO DI MILANO



## Motivating example

We have seen how to analyze very simple systems. But how can we analyze larger scenario, such as the process execution of an operative system supporting several concurrent tasks?



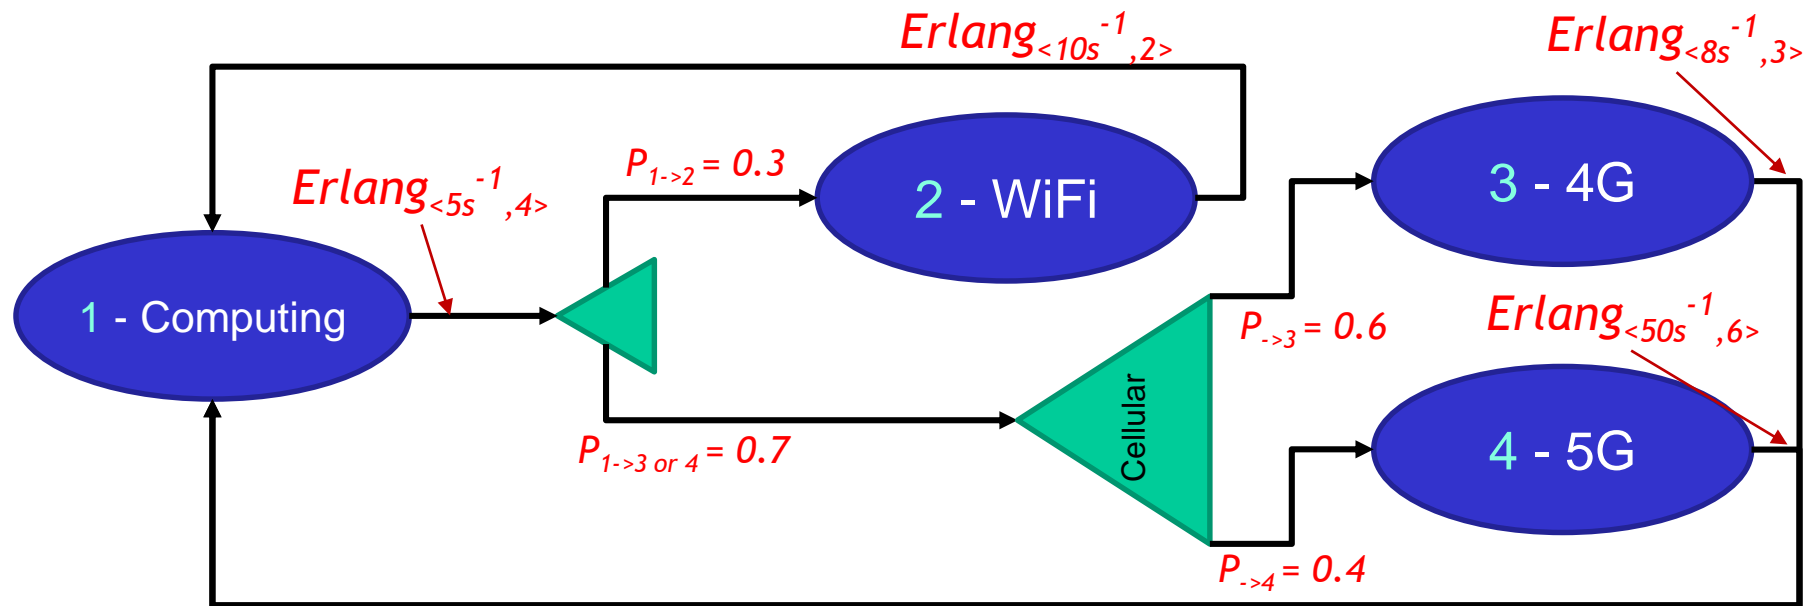


## Mobile Phone App (Choice)

Consider a Mobile Phone App that alternates between two phases: computing or sending data. The latter can be done using either WiFi or cellular (either 4G or 5G): depending on the technology this phase will take a different amount of time.

All timing have been measured and modelled with Erlang distributions with different number of stages and rate parameters.

Moreover, the probability of choosing either WiFi or cellular (4G or 5G) has been estimated.





## Mobile Phone App (Choice)

We would like to compute:

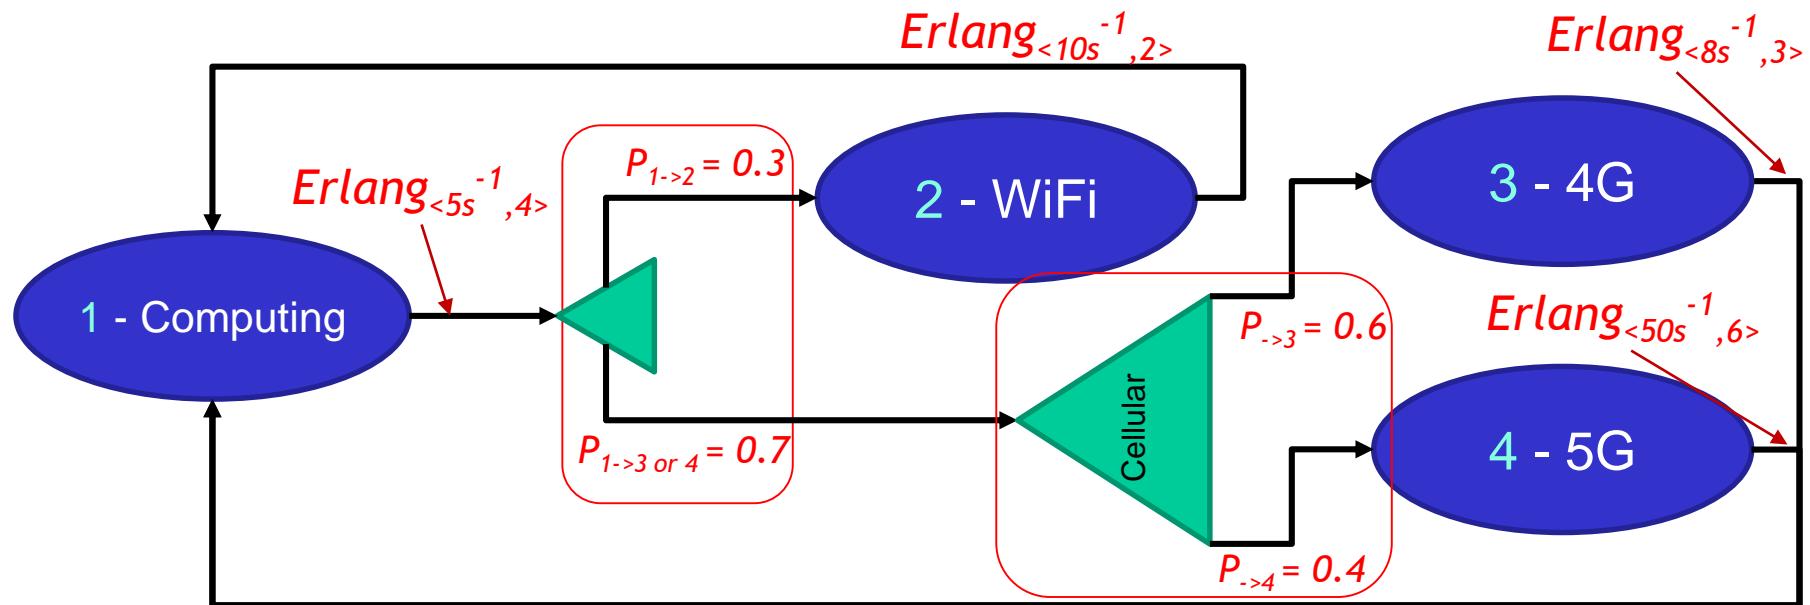
- State probabilities
- Throughput

In this case, the choice between the transition from state 1 to either state 2, 3 or 4, is guided by a discrete random variable.

$$P_{1 \rightarrow 2} = 0.3$$

$$P_{1 \rightarrow 3} = 0.7 \cdot 0.6 = 0.42$$

$$P_{1 \rightarrow 4} = 0.7 \cdot 0.4 = 0.28$$





## Mobile Phone App (Choice)

...

```
if s = 1 then
    u = rand()
    if u < 0.3
        ns = 2;
    else if u < 0.72
        ns = 3;
    else
        ns = 4;
    end
    dt = GenErlang(5,4);
end
if s = 2 then
    ns = 1;
    dt = GenErlang(10,2);
end
if s = 3 then
    ...
end
```

...

$$F_{1 \rightarrow 2} = 0.3$$

$$F_{1 \rightarrow 3} = 0.3 + 0.42 = 0.72$$

$$F_{1 \rightarrow 4} = 0.72 + 0.28 = 1$$



## Mobile Phone App (Choice)

```
Ts1 = Ts2 = Ts3 Ts4 = C = 0;  
...  
if s = 1 then  
    ...  
    dt = GenErlang(5,4);  
    Ts1 += dt;  
end  
if s = 2 then  
    ...  
    Ts2 += dt;  
    C ++;  
end  
if s = 3 then  
    ...  
    Ts3 += dt;  
    C ++;  
end  
...  
←
```

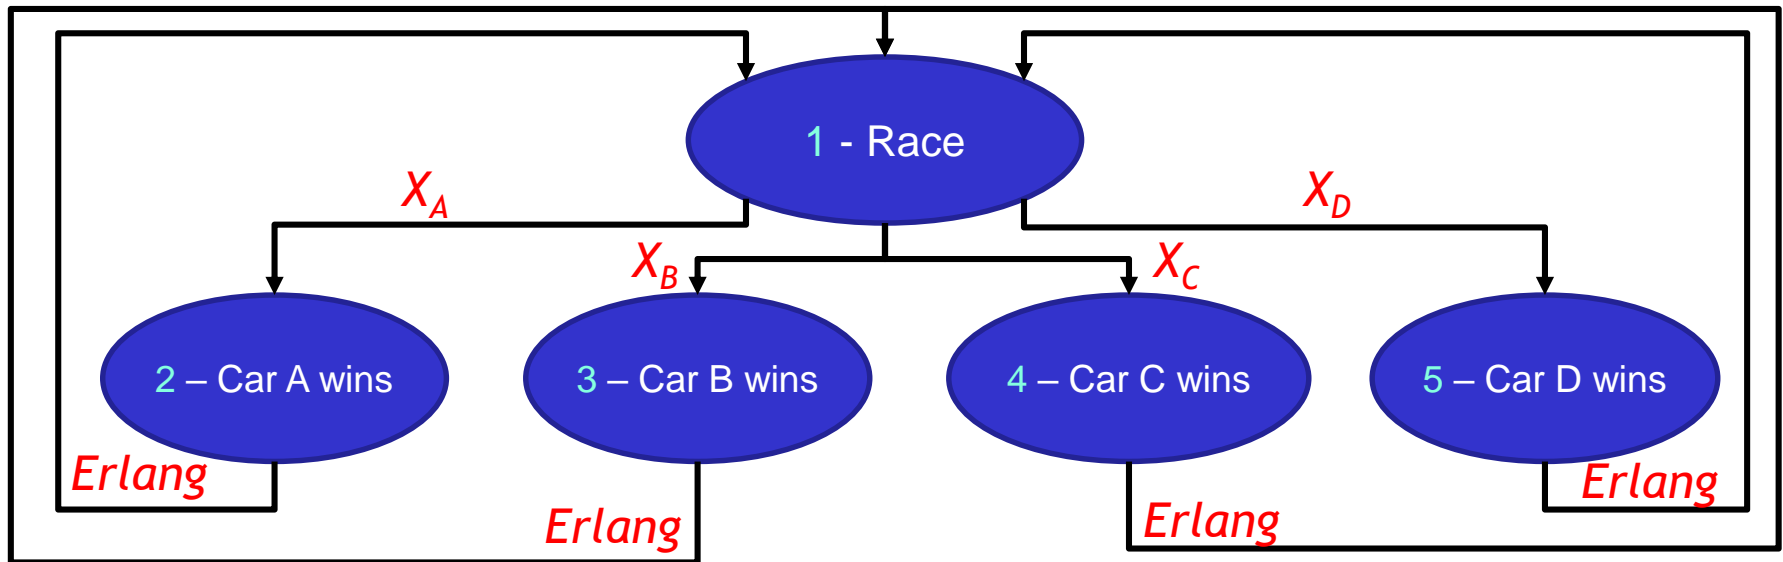
Also in this case,  
performance counters  
are updated in the code  
that handles each  
state.

```
Ps1 = Ts1 / T;  
Ps2 = Ts2 / T;  
Ps3 = Ts3 / T;  
Ps4 = Ts4 / T;  
X = C / T;
```



## Car Race Game

We want simulate a race, where four cars compete against each other. Each car has a different track completion time distribution. After a car wins the race, the game immediately stops and a new race restarts after an *Erlang* distributed amount of time. We can consider five states:

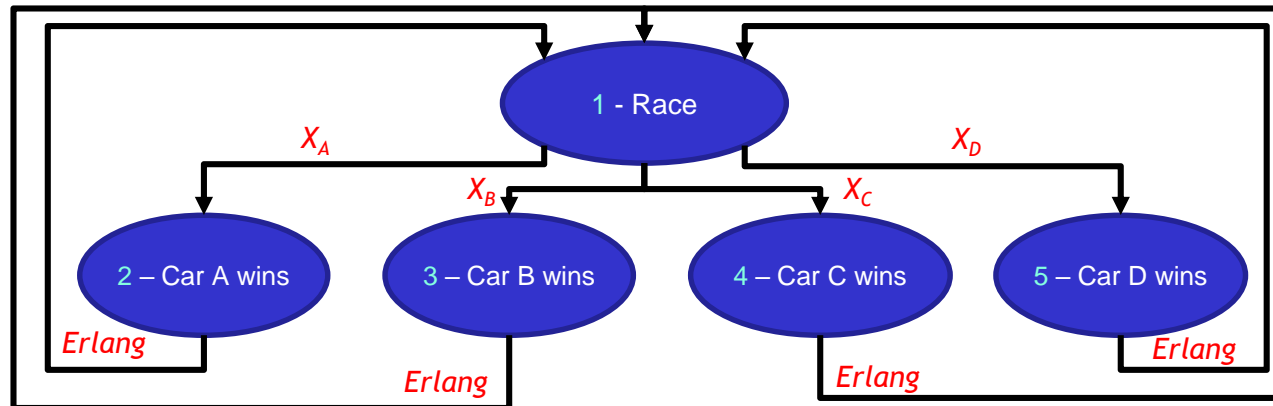




# Car Race Game

We would like to compute

- Victory probability for each car
- Average duration of a match.







...

```
if s = 1 then
```

```
    tA = GenXA();
```

```
    ...
```

```
    tD = GenXD();
```

```
    if (tA == min(tA, tB, tC, tD)) then
```

```
        ns = 2; dt = tA;
```

```
    end
```

```
    if (tB == min(tA, tB, tC, tD)) then
```

```
        ns = 3; dt = tB;
```

```
    ...
```

```
end
```

```
if s = 2 then
```

```
    ns = 1;
```

```
    dt = GenErlang();
```

```
end
```

```
if s = 3 then
```

```
    ns = 1;
```

```
    ...
```

...



...

```
if s = 1 then
```

```
    if (tA == min(tA, tB, tC, tD)) then
```

```
        ns = 2; dt = tA;
```

```
        WinA ++;
```

```
    end
```

```
    ...
```

```
    Races ++;
```

```
    RT(CurrRace,1) = dt;
```

```
    CurrRace ++;
```

```
end
```

...

```
PwinA = WinA / Races;
```

```
PwinB = WinB / Races;
```

```
PwinC = WinC / Races;
```

```
PwinD = WinD / Races;
```

```
AvgRT = mean(RT);
```

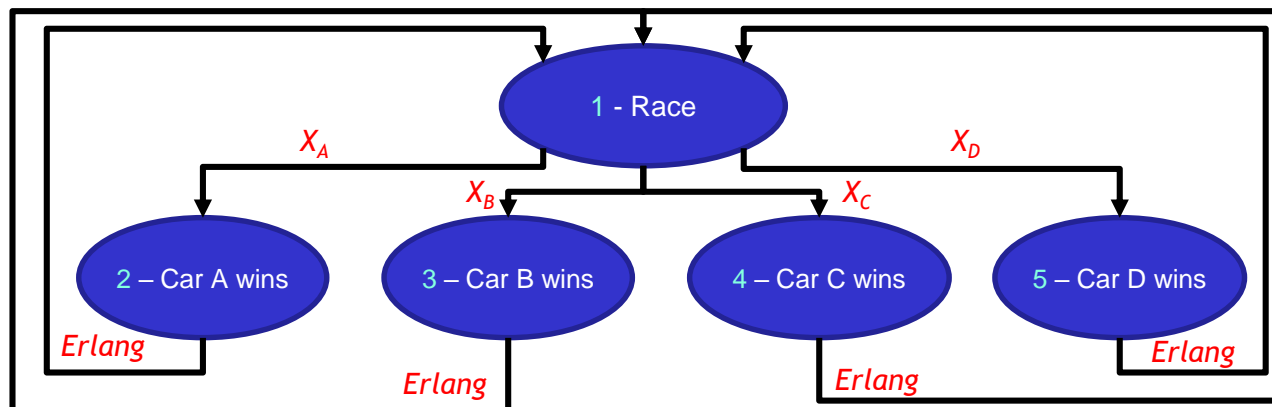


## Car Race Game

In this case, the choice of the next state is determined by the so called *Race policy*:

- the next state is chosen according to the path that has extracted the minimum time

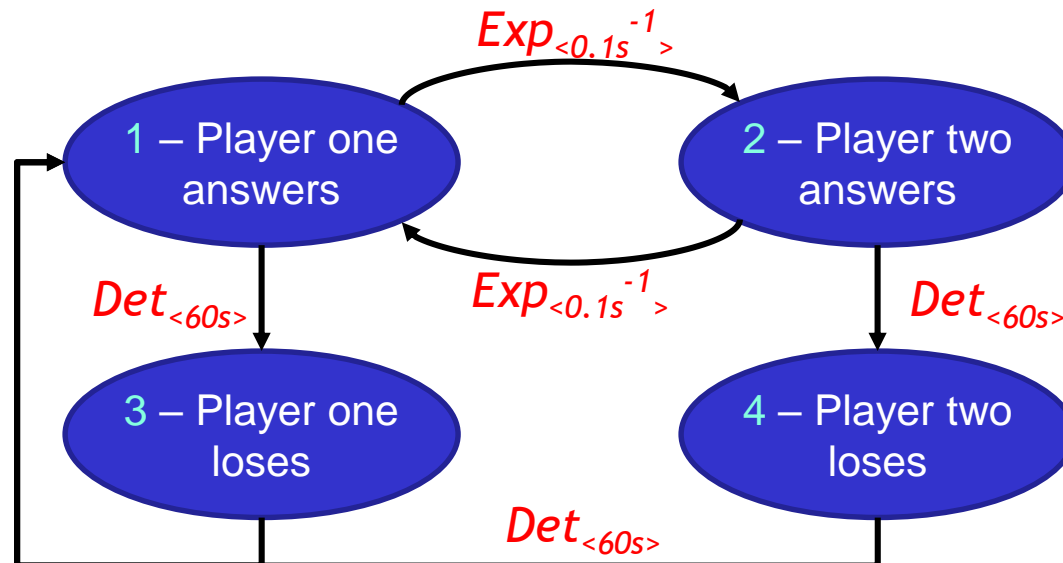
This is possible only in very few cases, in which all the transitions starts at the same time (such in this example), or when we have some special firing time distribution for the event.





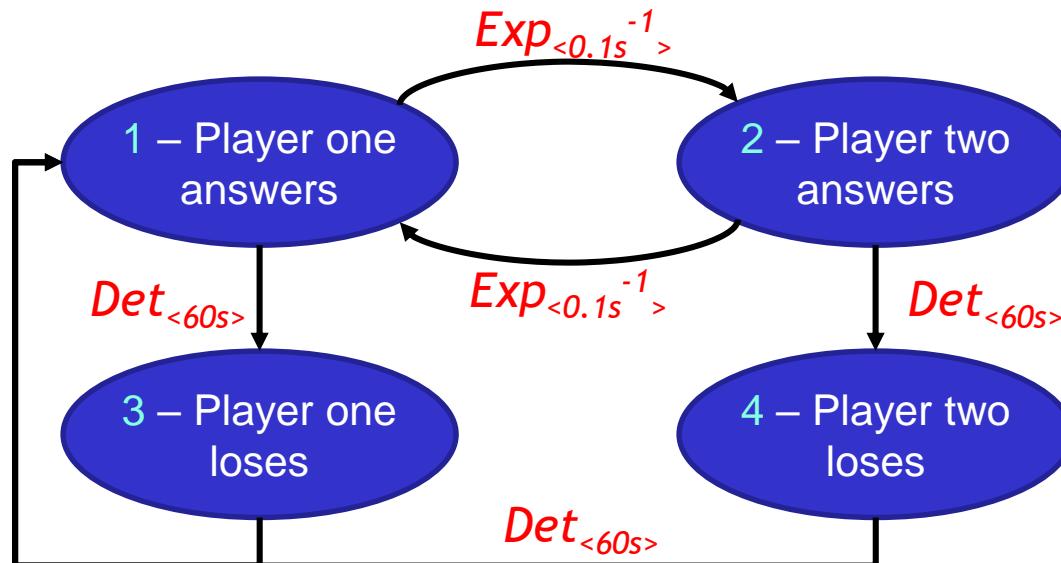
We want to simulate a two-contestant game: two people have 60 s. each, in which they have to answer a set of questions. Each time a player gives the correct answer, the turn passes to the other and hers timer stops. The player for which the timer expires first loses the game. Each player finds the correct answer in an exponentially distributed amount of time, with an average of 10 s.

Player one starts, and when a contestant loses, the game restarts after a deterministic time of 1 min.





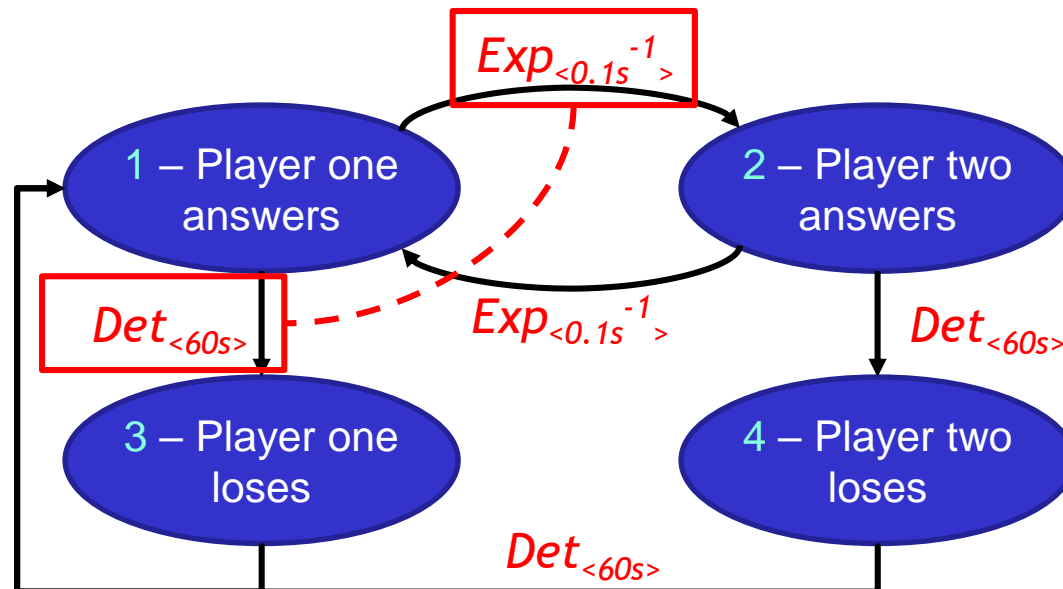
- Verify that the game is *fair*: each player has the same probability of winning.
- Compute the average length of one play.



In this case we cannot use the “race policy” shown in the previous example to handle states with more than an output arc:

- Race policy would result in each player always having one full minute to answer the question, without the time running out.

Timers and answers are *concurrent* event, and must be handled separately.





```
timerA = 60; timerB = 60;
...
  if s = 1 then
    tA = GenExp(0.1);
    if (tA < timerA) then
      ns = 2; dt = tA;
      timerA -= tA;
    else
      ns = 3; dt = timerA;
      timerA = 60;
    end
  end
  if s = 2 then
    ...
  end
  if s = 3 or s = 4 then
    ns = 1; dt = 60;
  ...
...
...
```



```
timerA = 60; timerB = 60; WinA = 0; WinB = 0; Match = 0;
```

```
...
```

```
  if s = 1 then
```

```
    tA = GenExp(0.1);
```

```
    if (tA < timerA) then
```

```
      ns = 2; dt = tA;
```

```
      timerA -= tA;
```

```
    else
```

```
      ns = 3; dt = timerA;
```

```
      timerA = 60; WinB ++;
```

```
    end
```

```
    RT(CurrMatch,1) += dt;
```

```
  end
```

```
  if s = 2 then
```

```
    ...
```

```
  end
```

```
  if s = 3 or s = 4 then
```

```
    ns = 1; dt = 60;
```

```
    CurrMatch ++; Match ++;
```


```
  ...
```

```
...
```

```
PWinA = WinA / Match;
```

```
PWinB = WinB / Match;
```

```
AvgRT = mean(RT);
```


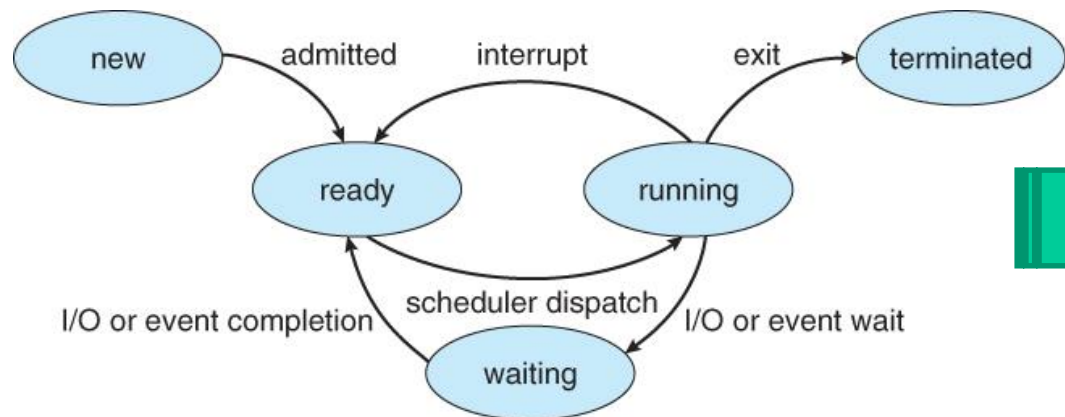






*Concurrency* can occur in many forms, and in general it cannot be handled with simple solutions like the ones previously seen.



Consider for example the state machine for the operative system process that we saw last time, applied to a batch process (a *closed model*): whenever a job ends, it restarts after a “think time”.





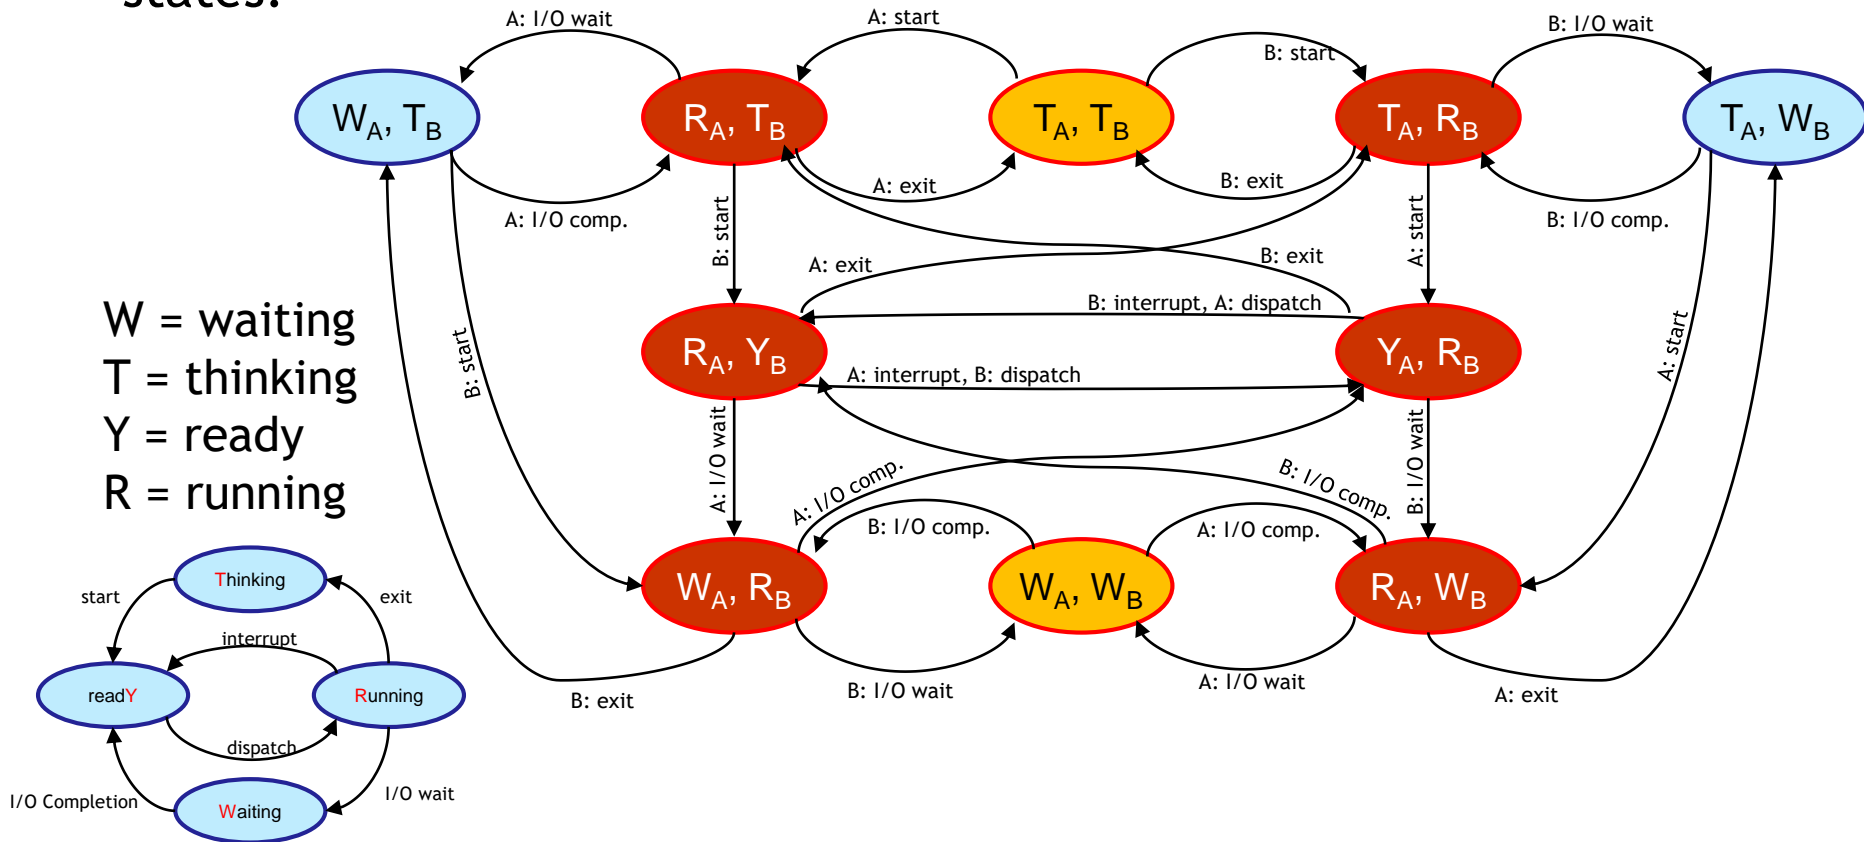
In case of a single batch process, jobs can be dispatched immediately, and the system would be in the “ready” state for a negligible amount of time. Such state could be removed.

Selection between “I/O wait” and “exit” in the running state can be performed using choice: a distribution of running before an event, and couple of probabilities for either ending or doing I/O.





In case of more jobs, different types of concurrency occur in almost all the states, and *choice* cannot be used to solve the problem in all states.



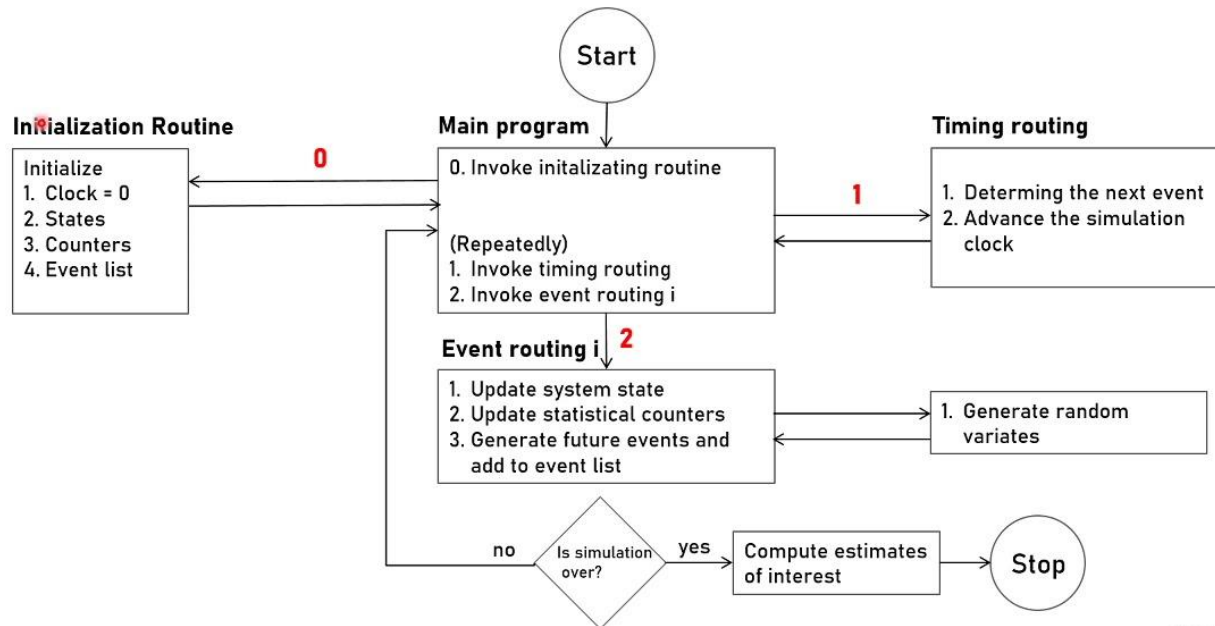


# Discrete Event Simulation

Concurrency in general can be handled using *Discrete Event Simulation*.

A complete description on how to perform Discrete Event Simulation, is however outside the scope of this course.

## Discrete Event Simulation Diagram Flow



Source: Law



## Exponential distribution and concurrency

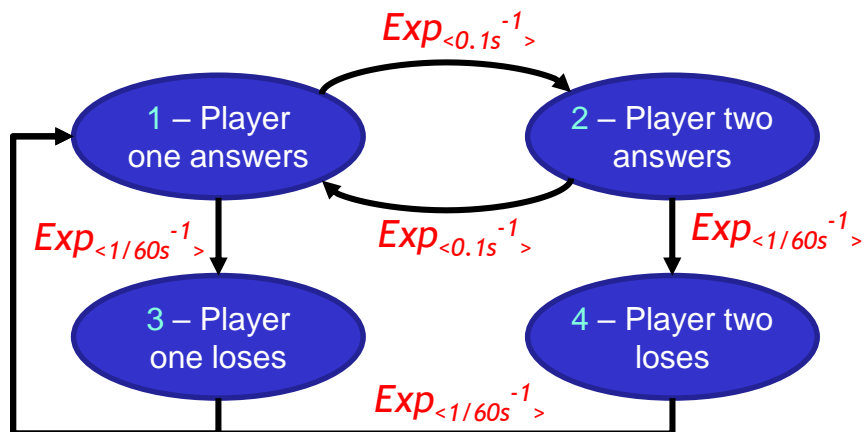
When ALL random timing follows an exponential distribution, analyzing systems become simpler thanks to its memory less property:

- Race policy becomes identical to concurrency
- Choice can be implemented with race policy, by embedding selection probability into the rates.



## Exponential distribution and concurrency

Race policy becomes identical to concurrency: for example, if in the “l'Eredità” game the timer would be exponentially distributed with an average of one minute, it could be analyzed in the following way:



```
if s = 1 then
    tA = GenExp(0.1);
    timer = GenExp(1/60);
    if (tA == min(tA, timer)) then
        ns = 2; dt = tA;
    else
        ns = 3; dt = timer;
    end
    ...
end
if s = 2 then
    ...
end
if s = 3 then
    ns = 1; dt = GenExp(1/60);
    ...
end
```

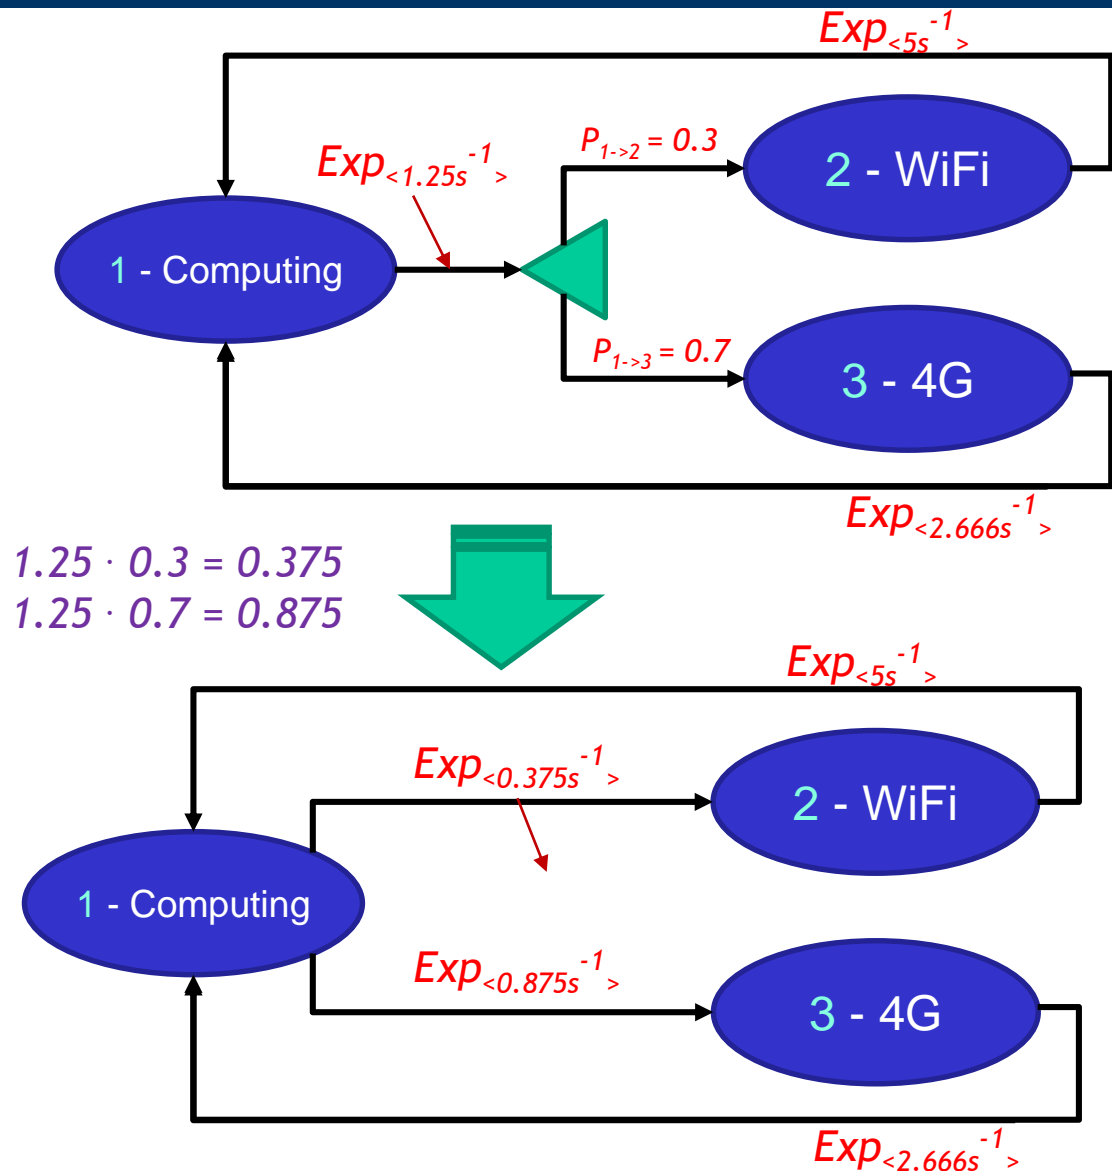
...



## Exponential distribution and concurrency

Choice can be implemented with race policy, using the product of the original exponential rate parameter with the selection probability.

For example, a 4G only version of the remote app model become:





## Analysis of Motivating Example

If we can assume exponential distributed timing, we can create a state machine that includes all the possible evolutions, and analyze it using race policy.

W = waiting  
T = thinking  
Y = ready  
R = running

